

# Volve G&G Dataset

The "Volve" oil field is an offshore oil and gas field located in the North Sea, approximately 190 kilometers west of Stavanger, Norway. It was discovered in 1993 and production began in 2008. The field is owned and operated by Equinor (formerly known as Statoil), which is one of the largest energy companies in the world. The field has estimated recoverable reserves of around 190 million barrels of oil equivalent, and produces mainly crude oil. Volve is a relatively small field compared to some of the other offshore fields in the North Sea, but it is still an important asset for Equinor and for Norway's overall oil and gas industry.

```
In [1]: from IPython.display import Image  
Image(url="https://www.norskipetroleum.no/factpages/3420717.jpg", width=700)
```

Out[1]:



## What kind of well-logs we will use?

- **NPHI** (neutron porosity): It is used to estimate the amount of fluids (usually hydrocarbons) contained in the rock formation by measuring the amount of neutron radiation that is emitted by the rock and reflected back to the sensor. This is important in the oil and gas industry to determine the potential productivity of a reservoir.
- **RHOB** (bulk density): It is used to determine the weight of the rock formation per unit volume. This is important for calculating the overall density of the rock formation and understanding its mechanical properties.

- **GR** (gamma ray): It is used to measure the amount of natural radiation that is emitted by the rock formation. This information can be used to identify certain types of rock formations, such as shale, and to estimate the amount of organic matter present in the formation.
- **RT** (resistivity): It is used to measure the electrical resistance of the rock formation to the flow of electric current. This information can be used to determine the presence and quality of fluids within the rock formation.
- **PEF** (photoelectric factor): It is used to measure the amount of X-ray radiation that is absorbed by the rock formation. This information can be used to identify certain types of rock formations, such as sandstone, and to estimate the amount of organic matter present in the formation.
- **CALI** (caliper): It is used to measure the diameter of the borehole. This information is important for determining the correct size of tools to be used for further measurements and for ensuring the stability of the borehole.
- **DT** (compressional travel time): It is used to measure the time it takes for a compressional (P-wave) sound wave to travel a known distance through the rock formation. This information can be used to determine the rock formation's mechanical properties, such as its porosity and permeability.

## Step 1 - importing data from Excel file and plotting the logs

```
In [2]: import pandas as pd
```

### Task 1

Read Excel file into a pandas dataframe and save it to the variable df using function pd.read\_excel

example: variable = pd.read\_excel('file\_name.xlsx')

print df - what is the index of your data?

```
In [3]: df = pd.read_excel('well_subset.xlsx')
display(df.head(10))
```

```
print("Index of the data is:", df.index)
```

	DEPTH	NPHI	RHOB	GR	RT	PEF	CALI	DT
0	2800.0	0.1425	2.4629	3.2562	1.7704	8.0126	8.5782	68.2803
1	2800.1	0.1416	2.4583	3.2575	1.7734	8.0124	8.6250	68.4759
2	2800.2	0.1436	2.4548	2.8439	1.8059	8.0316	8.6250	68.6693
3	2800.3	0.1454	2.4504	2.4479	1.8467	8.0325	8.6249	68.7748
4	2800.4	0.1509	2.4438	3.0292	1.9006	7.9983	8.5781	68.8805
5	2800.5	0.1549	2.4343	2.9266	1.9117	7.9443	8.5782	68.9858
6	2800.6	0.1573	2.4217	3.4017	1.8806	7.9051	8.6250	69.0042
7	2800.7	0.1632	2.4096	3.7842	1.8404	7.9249	8.6250	69.0204
8	2800.8	0.1679	2.4020	3.1949	1.8093	7.9677	8.6250	69.0371
9	2800.9	0.1703	2.4000	2.6821	1.7874	7.9796	8.6250	69.2040

Index of the data is: RangeIndex(start=0, stop=8001, step=1)

## Task 2

Set column DEPTH as the index using the following syntax

```
variable = variable.set_index('column_name')
```

Print the first 10 rows using the syntax

```
variable.head(number of rows)
```

```
In [4]: df = df.set_index('DEPTH')
df.head(10)
```

Out[4]:

	NPHI	RHOB	GR	RT	PEF	CALI	DT
DEPTH							
2800.0	0.1425	2.4629	3.2562	1.7704	8.0126	8.5782	68.2803
2800.1	0.1416	2.4583	3.2575	1.7734	8.0124	8.6250	68.4759
2800.2	0.1436	2.4548	2.8439	1.8059	8.0316	8.6250	68.6693
2800.3	0.1454	2.4504	2.4479	1.8467	8.0325	8.6249	68.7748
2800.4	0.1509	2.4438	3.0292	1.9006	7.9983	8.5781	68.8805
2800.5	0.1549	2.4343	2.9266	1.9117	7.9443	8.5782	68.9858
2800.6	0.1573	2.4217	3.4017	1.8806	7.9051	8.6250	69.0042
2800.7	0.1632	2.4096	3.7842	1.8404	7.9249	8.6250	69.0204
2800.8	0.1679	2.4020	3.1949	1.8093	7.9677	8.6250	69.0371
2800.9	0.1703	2.4000	2.6821	1.7874	7.9796	8.6250	69.2040

**Task 3** import library for plotting or install if needed. Specify the list of colors and plot the logs.

In [5]: `import matplotlib.pyplot as plt`

```

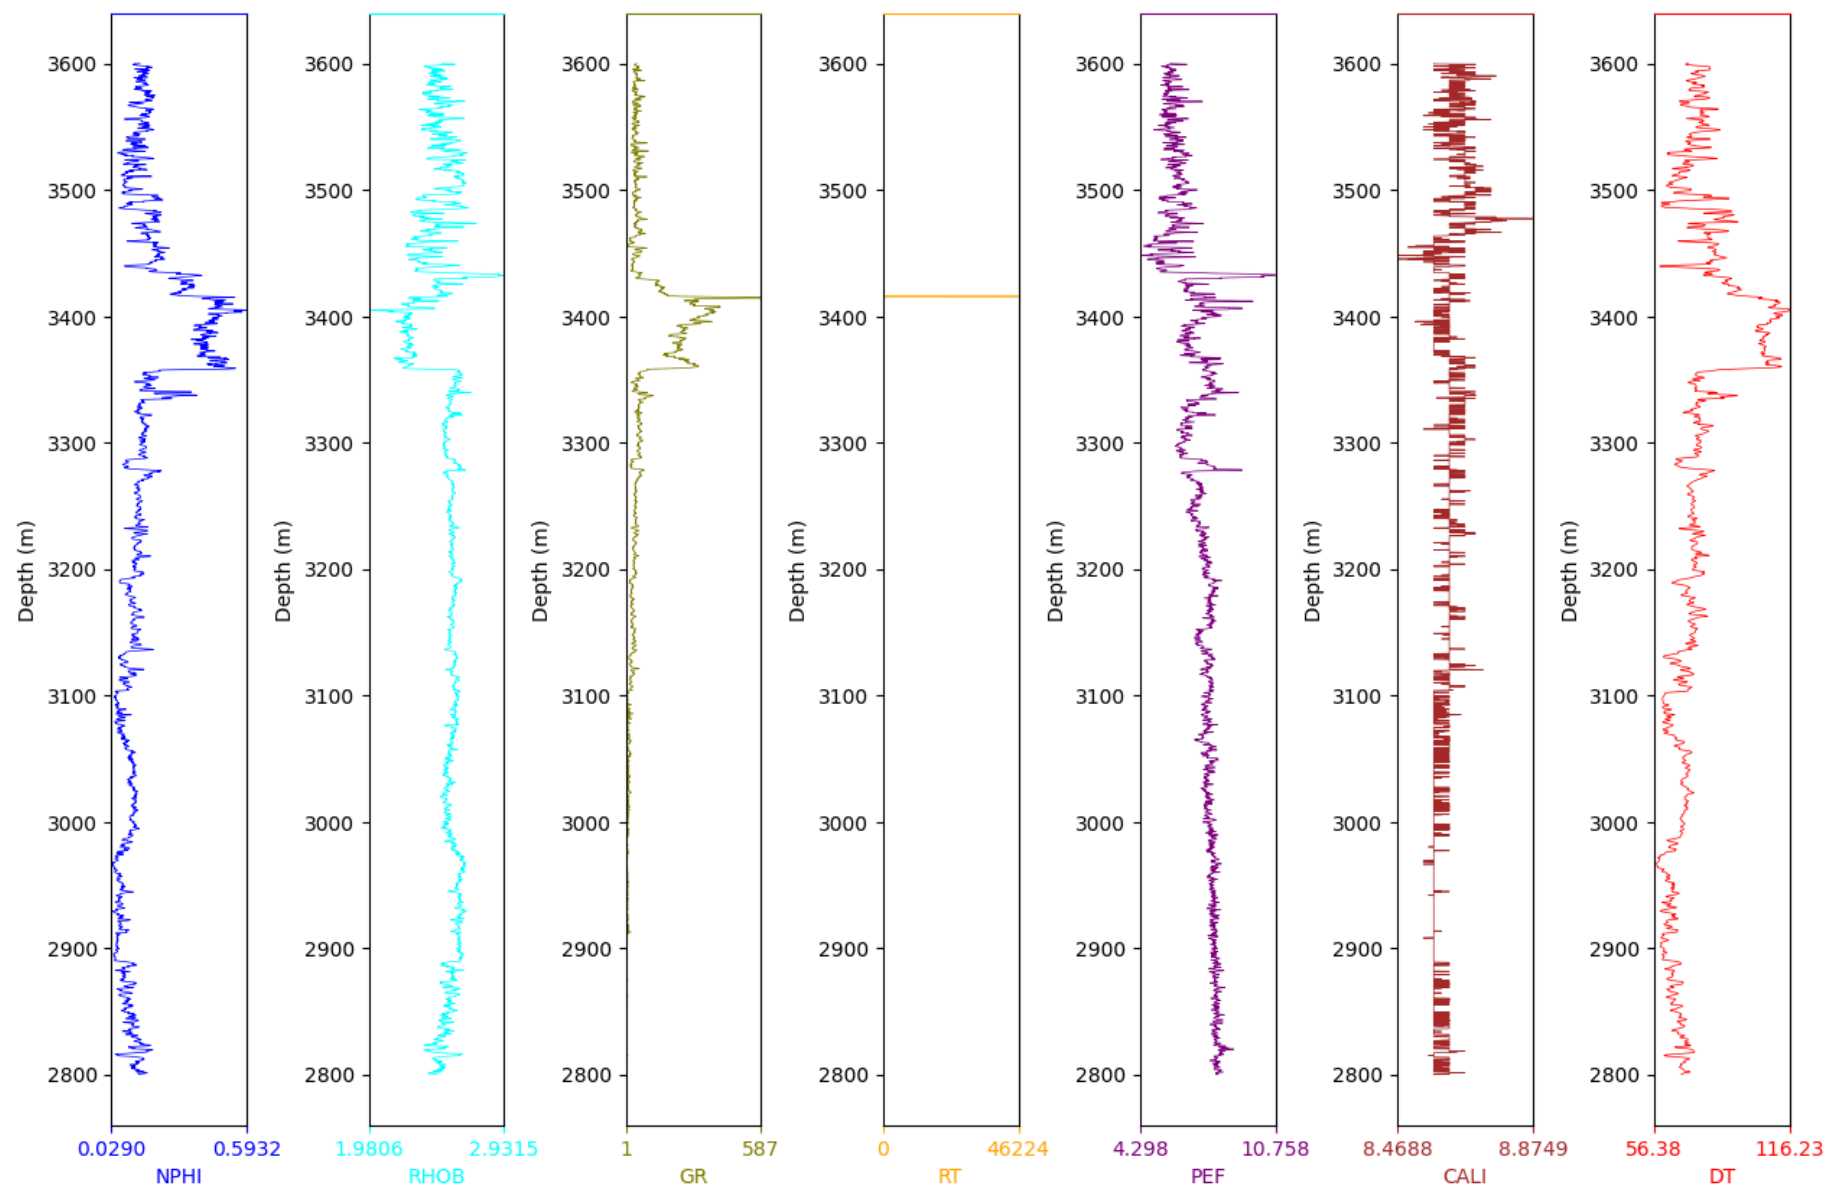
# Set up the plot axes using the number of columns for this put ncols value len(df.columns), set figsize to 15x10
fig, axs = plt.subplots(ncols=len(df.columns), figsize=(15, 10), gridspec_kw=dict(wspace=0.9))

# Define a list of 7 colors - see website https://matplotlib.org/stable/gallery/color/named_colors.html for ccolor list
colors = ['blue', 'cyan', 'olive', 'orange', 'purple', 'brown', 'red']

# Write a loop over the all columns in dataframe, specify linewidth to 0.5
for i, col in enumerate(df.columns):
    axs[i].plot(df.iloc[:,i], df.index, color=colors[i], linewidth=0.5)
    axs[i].set_xlabel(col)
    axs[i].xaxis.label.set_color(colors[i])

```

```
axs[i].set_xlim(df.iloc[:,i].min(), df.iloc[:,i].max())
axs[i].set_ylabel("Depth (m)")
axs[i].tick_params(axis='x', colors=colors[i])
axs[i].spines["top"].set_edgecolor(colors[i])
axs[i].title.set_color(colors[i])
axs[i].set_xticks([df.iloc[:,i].min(), df.iloc[:,i].max()])
```



**Task 4 Can you notice something strange? What with RT curve? write your comment**




The RT (Resistivity) curve (plotted in orange) appears as a flat horizontal line across the entire depth range. This suggests that all or most values are the same, or possibly that the curve contains extreme outliers or NaNs (missing values) that were not properly handled before plotting.

## Task 5 Check outliers - what is the outlier??

```
In [6]: # Set up the plot axes similar as before but use the fig size 20,5
fig, axs = plt.subplots(ncols=len(df.columns), figsize=(20,5), gridspec_kw=dict(wspace=0.5))

# write a loop function similar as above

for i, col in enumerate(df.columns):
    axs[i].boxplot(df[col])
    axs[i].set_xlabel(col)
    axs[i].set_ylabel("Value")
    axs[i].tick_params(axis='x', colors=colors[i])
    axs[i].spines["top"].set_edgecolor(colors[i])
    axs[i].title.set_color(colors[i])
```



We need to remove weird high values of RT, but what to do with this observations? We can replace them using interpolation from the nearest samples



## Task 6 Correct RT curve

```
In [7]: import numpy as np

# Replace RT Log values greater than 100 with NaN
df.loc[df['RT'] > 100, 'RT'] = np.nan

# Interpolate NaN values based on nearest values
df['RT'] = df['RT'].interpolate(method='nearest')

# Print the resulting DataFrame
print(df)
```

	NPHI	RHOB	GR	RT	PEF	CALI	DT
DEPTH							
2800.0	0.1425	2.4629	3.2562	1.7704	8.0126	8.5782	68.2803
2800.1	0.1416	2.4583	3.2575	1.7734	8.0124	8.6250	68.4759
2800.2	0.1436	2.4548	2.8439	1.8059	8.0316	8.6250	68.6693
2800.3	0.1454	2.4504	2.4479	1.8467	8.0325	8.6249	68.7748
2800.4	0.1509	2.4438	3.0292	1.9006	7.9983	8.5781	68.8805
...	...	...	...	...	...	...	...
3599.6	0.1289	2.5771	44.3674	2.3147	6.1787	8.5781	70.1850
3599.7	0.1259	2.5490	43.5794	2.3004	5.9839	8.5781	70.3162
3599.8	0.1312	2.5246	44.6774	2.2336	5.8875	8.5781	70.5137
3599.9	0.1365	2.5003	45.4844	2.1827	5.7913	8.5781	70.7711
3600.0	0.1470	2.4950	47.8596	2.1170	5.7226	8.5784	71.3462

[8001 rows x 7 columns]

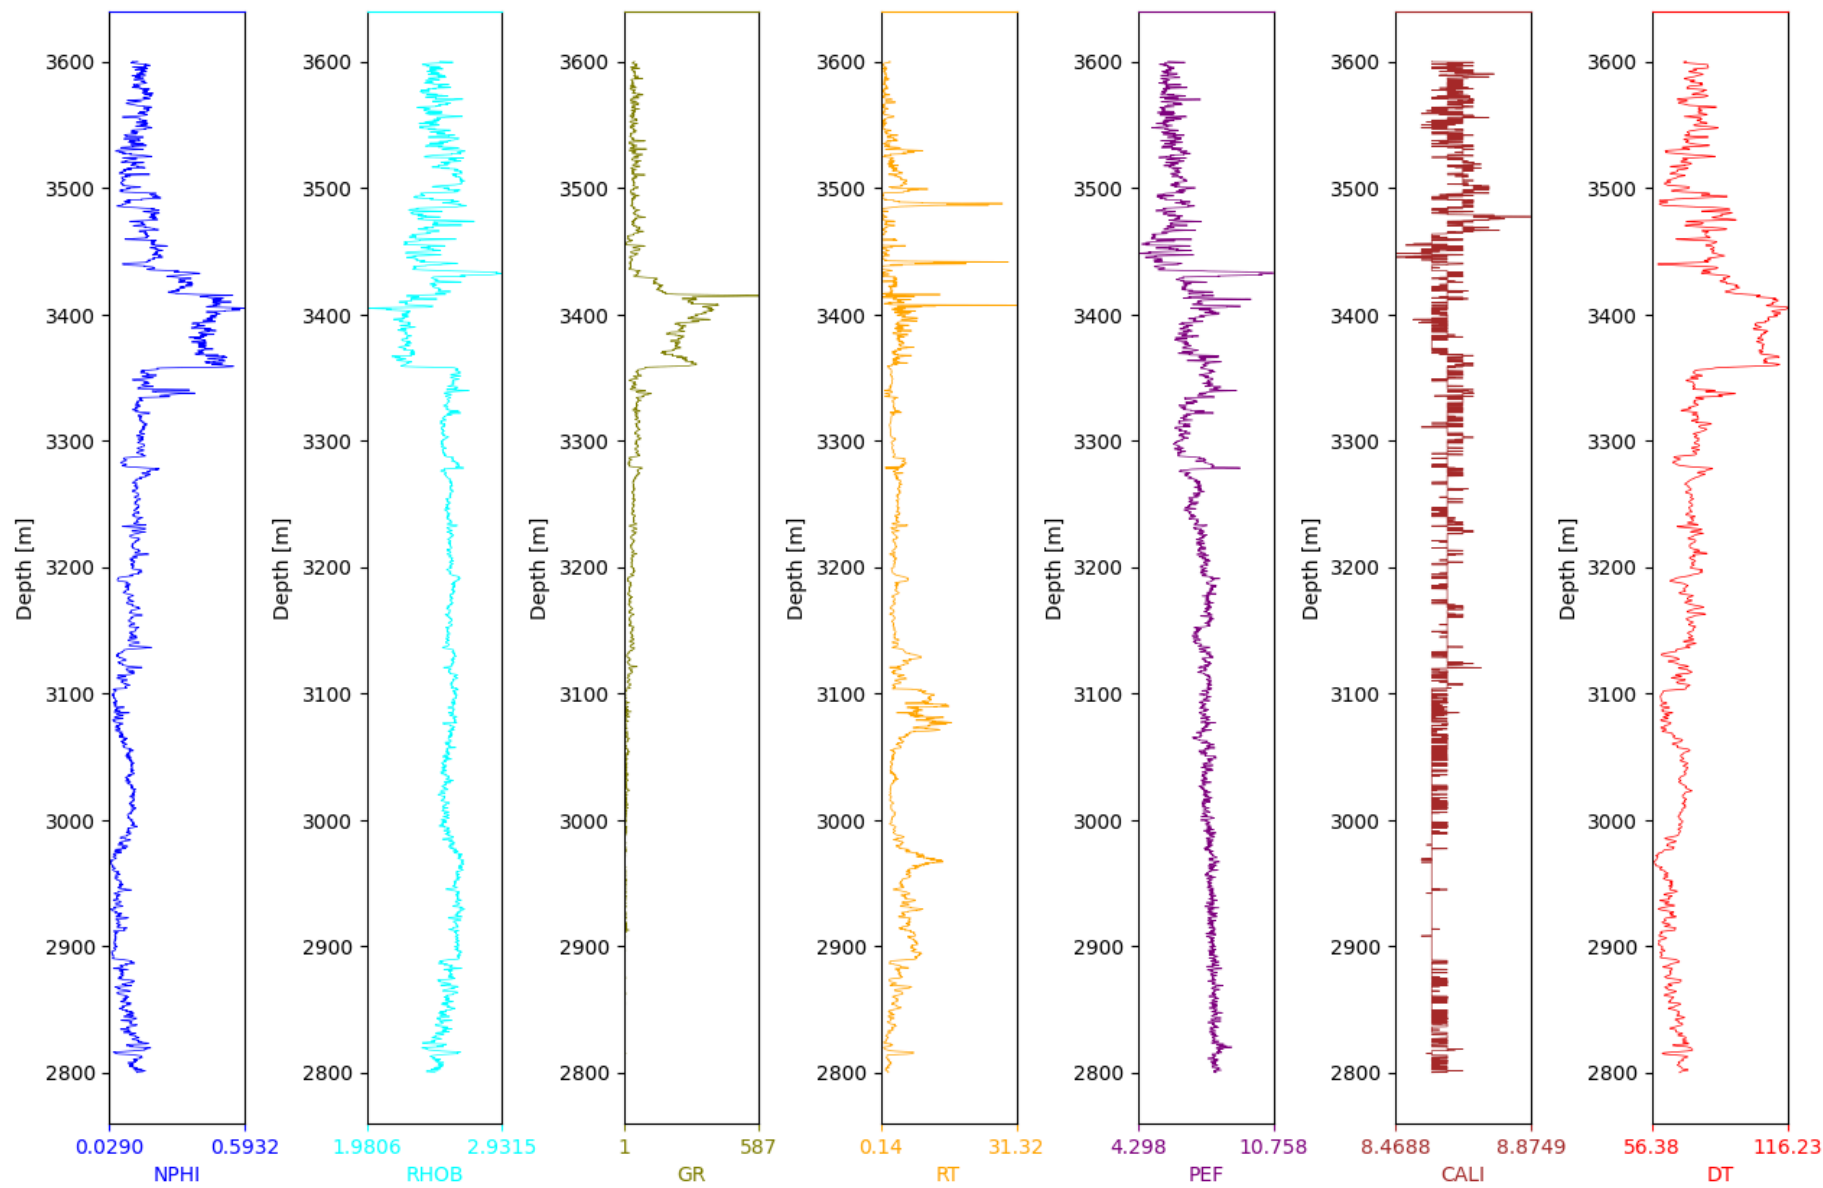
See if it helped! We will visualize data again

```
In [8]: import matplotlib.pyplot as plt

# Set up the plot axes using the number of columns for this put ncols value len(df.columns), set figsize to 15x10
fig, axs = plt.subplots(ncols=len(df.columns), figsize=(15, 10), gridspec_kw=dict(wspace=0.9))

# Define a List of 7 colors - see website https://matplotlib.org/stable/gallery/color/named_colors.html for ccolor list
colors = ['blue', 'cyan', 'olive', 'orange', 'purple', 'brown', 'red']
```

```
# Write a loop over the all columns in dataframe, specify linewidth to 0.5
for i, col in enumerate(df.columns):
    axs[i].plot(df.iloc[:,i], df.index, color=colors[i], linewidth=0.5)
    axs[i].set_xlabel(col)
    axs[i].xaxis.label.set_color(colors[i])
    axs[i].set_xlim(df.iloc[:,i].min(), df.iloc[:,i].max())
    axs[i].set_ylabel("Depth [m]")
    axs[i].tick_params(axis='x', colors=colors[i])
    axs[i].spines["top"].set_edgecolor(colors[i])
    axs[i].title.set_color(colors[i])
    axs[i].set_xticks([df.iloc[:,i].min(), df.iloc[:,i].max()])
```




Did it help?

Yes - after removing and interpolating the extreme RT values (those above 100), the RT curve now appears more continuous and realistic, without the large vertical jumps or flat lines previously caused by outliers. This improves both the visual clarity and geological interpretability of the log data.

```
In [9]: # Set up the plot axes similar as before but use the fig size 20,5
fig, axs = plt.subplots(ncols=len(df.columns), figsize=(20,5), gridspec_kw=dict(wspace=0.5))

# write a loop function similar as above

for i, col in enumerate(df.columns):
    axs[i].boxplot(df[col])
    axs[i].set_xlabel(col)
    axs[i].set_ylabel("Value")
    axs[i].tick_params(axis='x', colors=colors[i])
    axs[i].spines["top"].set_edgecolor(colors[i])
    axs[i].title.set_color(colors[i])
```






## Task 7 check the outliers visualization using seaborn library



check the seaborn website: <https://seaborn.pydata.org/examples/index.html> write in the comment 3 visualizations that you can use in your daily tasks

```
In [10]: import seaborn as sns



for i, col in enumerate(df.columns):
    fig, ax = plt.subplots(figsize=(4, 4))
    sns.boxplot(y=df[col], color=colors[i], orient='v')
```











## Comment:

Three visualizations from seaborn useful in daily tasks:

1. Boxplot – for detecting outliers and understanding data spread.
2. Heatmap – for visualizing correlation matrices and identifying relationships.
3. Pairplot – for exploring pairwise relationships and distributions of multiple variables.

## Task 8 - Remove outliers that are 3 standard deviations from the mean in window of 100 samples

```
In [11]: # set window size to 100
window_size = 100

# calculate z-scores for each column using rolling window
z_scores = (df - df.rolling(window=window_size).mean()) / df.rolling(window=window_size).std()

# remove rows where any z-score is greater than 3
df_noout = df[(np.abs(z_scores) < 3).all(axis=1)]

# print cleaned dataframe
print(df_noout)
```

	NPHI	RHOB	GR	RT	PEF	CALI	DT
DEPTH							
2809.9	0.1427	2.4589	4.2370	2.0991	7.9082	8.6249	70.0120
2810.0	0.1402	2.4615	4.1468	2.0325	7.9031	8.5782	69.9705
2810.1	0.1393	2.4696	4.2389	1.9359	7.9372	8.6015	69.8849
2810.2	0.1348	2.4811	3.6711	1.8266	7.9944	8.5781	69.7992
2810.3	0.1309	2.4924	2.8638	1.7334	8.0531	8.5781	69.7205
...	...	...	...	...	...	...	...
3599.0	0.1207	2.4931	41.2866	1.4664	5.8496	8.7030	71.5955
3599.1	0.1200	2.5050	42.4855	1.6046	6.0435	8.6718	71.1964
3599.2	0.1240	2.5236	42.9226	1.7066	6.2858	8.6252	70.8162
3599.9	0.1365	2.5003	45.4844	2.1827	5.7913	8.5781	70.7711
3600.0	0.1470	2.4950	47.8596	2.1170	5.7226	8.5784	71.3462

[6618 rows x 7 columns]

```
In [12]: # Set up the plot axes
fig, axs = plt.subplots(ncols=len(df.columns), figsize=(20,5), gridspec_kw=dict(wspace=0.5))


# make a for loop for visualization using df dataset

for i, col in enumerate(df.columns):
    sns.boxplot(y=df[col], ax=axs[i], color='lightblue')
    axs[i].set_title(f'Original {col}')

# Set up the plot axes
fig, axs = plt.subplots(ncols=len(df.columns), figsize=(20,5), gridspec_kw=dict(wspace=0.5))

# make a for loop for visualization using df_noout dataset

for i, col in enumerate(df_noout.columns):
    sns.boxplot(y=df_noout[col], ax=axs[i], color='lightgreen')
    axs[i].set_title(f'Cleaned {col}')
```



## Step 2 - Exploratory Data Analysis

EDA (Exploratory Data Analysis) is the process of analyzing and visualizing data in order to extract insights, patterns, and trends. It is typically one of the first steps in data analysis and is used to gain a better understanding of the data and its characteristics. EDA can help identify outliers, missing values, and any other issues with the data that may need to be addressed before further analysis. It can also help in selecting appropriate statistical methods and models for data analysis. EDA involves using a range of techniques such as histograms, scatter plots, box

plots, and correlation matrices to explore the data visually and identify relationships between variables. EDA is an important part of data science and plays a crucial role in the data analysis process.

## Histograms

Histograms are graphical representations of the distribution of data. They display the frequency distribution of a variable by creating a set of contiguous and non-overlapping intervals (or bins) along the range of the variable and then plotting the count or proportion of observations that fall within each bin. By examining a histogram, one can see the central tendency, variability, and shape of the distribution of the data. Additionally, it can also help identify any outliers or unusual patterns in the data. Overall, histograms are useful for understanding the distribution of a variable and gaining insight into the underlying patterns and characteristics of the data.

### Optimal bin size

**Freedman-Diaconis rule** is a method for determining the bin width of a histogram in statistical data analysis. The rule uses the interquartile range (IQR) of the data and the total number of samples to calculate an appropriate bin width. The bin width is important because it determines the smoothness of the histogram and can affect the interpretation of the data. The Freedman-Diaconis rule aims to create a histogram with sufficient detail to reveal the underlying distribution of the data while avoiding oversmoothing or undersmoothing. It is considered a robust method for determining bin width because it is less sensitive to outliers than other methods such as Sturges' rule or Scott's rule. The Freedman-Diaconis rule has been widely adopted in various fields, including economics, environmental science, and medical research.

```
In [13]: # Set up the plot axes as before, use size 20,5
fig, axs = plt.subplots(ncols=len(df.columns), figsize=(20, 5), gridspec_kw=dict(wspace=0.5))

#write a loop
for i, col in enumerate(df_noout.columns):

    # Calculate the bin size using the Freedman-Diaconis rule
    # specify the quartiles
    q75, q25 = np.percentile(df_noout[col], [75, 25])
    # calculate iqr (interquartile range)
    iqr = q75 - q25


    # calculate the width as h equals to double iqr divided by the cube root of the number of observation -> n = len(df[col])
```

```

h = 2 * iqr / (len(df[col]) ** (1/3))
bins = int((df_noout[col].max() - df_noout[col].min()) / h)

axs[i].hist(df_noout[col], bins=bins, color=colors[i])
axs[i].set_xlabel(col)
axs[i].set_ylabel("N")
axs[i].tick_params(axis='x', colors=colors[i])
axs[i].spines["top"].set_edgecolor(colors[i])
axs[i].title.set_color(colors[i])

```



## Pairplot

A pairplot is a graphical tool used in data analysis and visualization that creates pairwise scatterplots and histograms for a given dataset. It is a useful method for exploring the relationship between multiple variables in a dataset. Each scatterplot in the pairplot shows the relationship between two variables in the dataset, while the histograms show the distribution of each variable individually.

By examining the pairplot, we can gain insights into the relationships between variables in the dataset. We can identify variables that are strongly correlated, positively or negatively, as well as variables that are not correlated at all. We can also see the distribution of each variable, including whether they are normally distributed or skewed, and identify any outliers. Pairplots can be useful for identifying potential patterns or trends in the data, as well as for detecting any issues with the data, such as missing or erroneous values. Overall, pairplots are a useful tool for exploratory data analysis and for gaining a better understanding of the relationships between variables in a dataset.

```
In [14]: cols = ['NPHI', 'RHOB', 'GR', 'RT', 'PEF', 'CALI', 'DT']

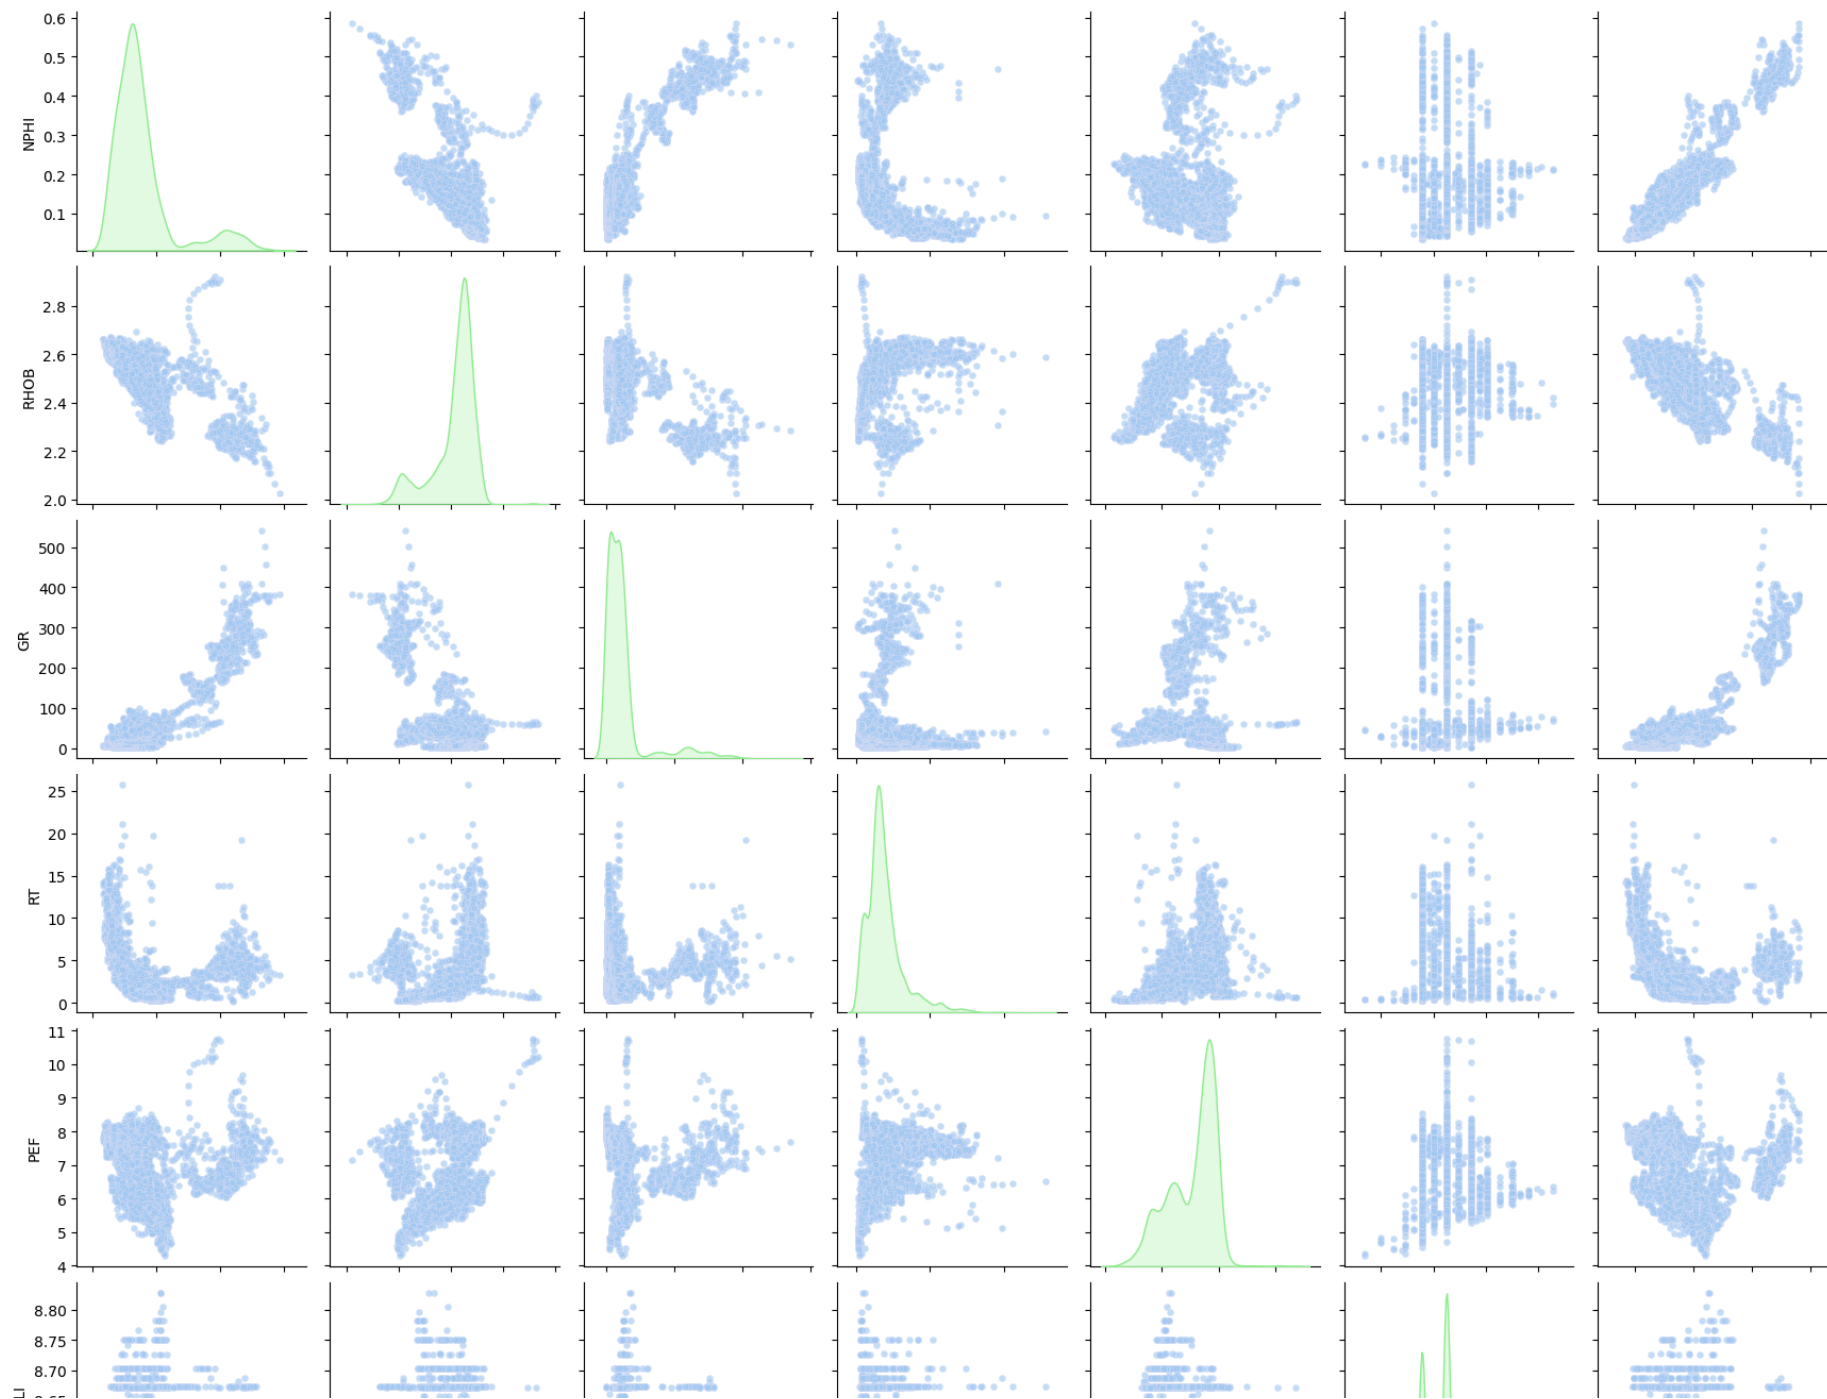
# Define colors for histograms and scatters
hist_color = 'lightgreen'
scatter_color = 'lavender'


# Set the color palette
sns.set_palette("pastel")

# Create pairplot with different colors for histograms and scatters put as vars your column list, and use diag_kind as 'kde',
sns.pairplot(df_noout, vars=cols, diag_kind='kde',
             plot_kws = {'alpha': 0.6, 's': 20, 'edgecolor': scatter_color},
             diag_kws = {'color': hist_color})
```

```
Out[14]: <seaborn.axisgrid.PairGrid at 0x2ca37e17770>
```







#### Observations from the Pairplot **Distributions (Diagonal KDE Plots):**

- NPHI (Neutron Porosity): Positively skewed with a peak around 0.15–0.2.
- RHOB (Bulk Density): Appears normally distributed, centered around 2.4–2.5 g/cc.
- GR (Gamma Ray): Right-skewed with a concentration below 100 API.
- RT (Resistivity): Highly right-skewed, with most values clustered below 10 ohm-m.
- PEF (Photoelectric Effect): Slightly right-skewed, peak around 2–3.
- CALI (Caliper): Distribution is spiky with evident discretization, suggesting it may be recorded at specific intervals.
- DT (Sonic Travel Time): Nearly normal distribution, centered around 90–100  $\mu\text{s}/\text{ft}$ .

#### Variable Relationships (Scatterplots):

- NPHI vs RHOB: Inversely correlated – common in petrophysics due to fluid/gas effects.
- GR vs RT: Suggests possible clustering, potentially indicating lithology or fluid zones.
- NPHI vs PEF: Shows non-linear trends, possibly hinting at lithology-based clusters.
- RHOB vs DT: Shows a curved relationship — consistent with compaction trends in sedimentary formations.

- CALL vs other variables: Appears more discretized and less informative visually.
- RT vs others: Exhibits a wide dynamic range; log scaling might help in further plots.

#### Potential Outliers:

- There are clear outliers in several scatterplots (e.g.,  $RT > 100$ ,  $GR > 300$ ), although the dataset is labeled as `df_noout` — further cleaning might be necessary.
- Some variables such as RT and GR show vertical clustering lines, which may indicate sensor saturation or digitization artifacts.

#### Clustering Insight:

- Several plots exhibit patterns that could be exploited for clustering (e.g., NPHI-RHOB, RHOB-DT, PEF-NPHI).
- Based on the visual groupings in scatterplots, clustering algorithms like KMeans, DBSCAN, or Gaussian Mixture Models (GMM) could be used to segment the data meaningfully.

#### Summary Observations from the Pairplot

- **Distributions:** Most variables (e.g., RT, GR, NPHI) are right-skewed; RHOB and DT appear more normally distributed.
- **Relationships:** Clear inverse correlation between NPHI and RHOB. Non-linear patterns observed between variables like RHOB-DT and PEF-NPHI.
- **Outliers:** Some extreme values are still present (e.g., high RT, GR), possibly due to measurement artifacts or residual noise.
- **Clustering Potential:** Visual clusters are noticeable in several plots (e.g., NPHI vs RHOB), suggesting that clustering (e.g., KMeans, DBSCAN, GMM) may reveal lithological or reservoir-related groupings.

## Correlation matrix

Correlation matrix is a table that displays the correlation coefficients between different variables in a dataset. It is commonly used in statistics and data analysis to identify patterns and relationships between variables.

The correlation coefficient is a statistical measure that ranges from -1 to 1, indicating the strength and direction of the linear relationship between two variables. A value of 1 indicates a perfect positive correlation, meaning that when one variable increases, the other variable increases proportionally. A value of -1 indicates a perfect negative correlation, meaning that when one variable increases, the other variable decreases proportionally. A value of 0 indicates no correlation, meaning that there is no relationship between the variables.

The correlation matrix is a square matrix where the diagonal contains the correlation coefficient between each variable and itself, which is always 1. The upper and lower triangles of the matrix contain the correlation coefficients between each pair of variables, with duplicates reflected across the diagonal. A correlation matrix can be visualized as a heat map, where the color of each cell represents the magnitude of the correlation coefficient. Correlation matrices are commonly used in data analysis, machine learning, and other applications to identify relationships between variables, detect multicollinearity, and perform feature selection.


```
In [15]: # Define the features to include in the correlation matrix
cols = ['NPHI', 'RHOB', 'GR', 'RT', 'PEF', 'CALI', 'DT']

# Calculate the correlation matrix using the following syntax dataframe[columns].corr()
corr_matrix = df_noout[cols].corr()

# Set the color palette
sns.set_palette("pastel")

# Create a heatmap of the correlation matrix, set annot to True
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
```

```
Out[15]: <Axes: >
```



### Correlation Matrix Observations

#### Strong positive correlations:

- NPHI and GR (0.92): Likely linked to shale content.
- NPHI and DT (0.96): Indicates similar response to porosity.
- GR and DT (0.90): Suggests GR increases with formation slowness.

#### Strong negative correlations:

- NPHI and RHOB (-0.79): Classic inverse relationship in porosity zones.

- RHOB and DT (-0.81): Denser formations have lower travel times.

**Moderate correlations:**

- RHOB and PEF (0.48): Both relate to formation lithology.
- RT and PEF (0.38): Slight lithological control over resistivity.

**Low or negligible correlations:**

- CALL shows weak correlation with all variables - it might carry limited clustering value.
- RT generally shows weak correlations with other logs, which is common due to its sensitivity to fluid content rather than rock matrix.

## Step 3 - data normalization

Data normalization, also known as feature scaling, is the process of transforming data into a common scale or range in order to facilitate data analysis and improve the performance of machine learning algorithms. Normalization is important because many machine learning algorithms are sensitive to the scale and distribution of input features, and may perform poorly or inaccurately if the features are not on a similar scale.

Normalization involves rescaling the features of a dataset to have a mean of 0 and a standard deviation of 1, or scaling the features to a range between 0 and 1. The normalization method used depends on the specific data and the requirements of the analysis or algorithm being used. Common methods of normalization include Min-Max scaling, Z-score normalization, and Log transformation.

Min-Max scaling involves scaling the features to a range between 0 and 1, where the minimum value of the feature is transformed to 0 and the maximum value is transformed to 1. Z-score normalization involves transforming the features so that they have a mean of 0 and a standard deviation of 1, which can be accomplished by subtracting the mean from each value and then dividing by the standard deviation. Log transformation is another normalization technique used for data that is highly skewed or has a wide range of values, and involves applying a logarithmic function to the data to transform it into a more normal distribution.

Overall, data normalization is an important preprocessing step in data analysis and machine learning, as it can help to improve the accuracy and performance of models and algorithms, reduce overfitting, and ensure that features are on a similar scale.

## 3.1 - Transform Resistivity Log to log scale

Resistivity data in well logs is typically measured in ohm-meters, and the values can span several orders of magnitude, making it difficult to visualize and analyze the data directly. To address this issue, resistivity data is often transformed using a logarithmic scale, which compresses the data into a more manageable range.

The logarithmic scale is a nonlinear scale that compresses large values into a smaller range, while expanding small values. This allows for a more accurate visualization of the data and makes it easier to identify patterns and trends. In particular, the use of logarithmic scales is useful for resistivity data because the range of resistivity values encountered in well logs can be very large, spanning several orders of magnitude.

```
In [16]: df_noout['RT'] = np.log10(df_noout['RT'])
```

C:\Users\Alexa\AppData\Local\Temp\ipykernel\_5184\3476924017.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_noout['RT'] = np.log10(df_noout['RT'])
```

```
In [17]: #check data
df_noout['RT']
```

```
Out[17]: DEPTH
2809.9    0.322033
2810.0    0.308031
2810.1    0.286883
2810.2    0.261643
2810.3    0.238899
...
3599.0    0.166252
3599.1    0.205367
3599.2    0.232132
3599.9    0.338994
3600.0    0.325721
Name: RT, Length: 6618, dtype: float64
```



## 3.2 - Transform data with skewed distribution


The power transform with Yeo-Johnson method is a data transformation technique used to normalize a dataset that has a skewed distribution. It is a variant of the Box-Cox transformation, which is used to normalize data that has a positive skew.

The Yeo-Johnson method is a modified version of the Box-Cox transformation that can be applied to both positively and negatively skewed data, as well as data that includes zero or negative values. It works by applying a power transformation to the data that varies based on the value of a lambda parameter, which is estimated from the data.

The Yeo-Johnson method is implemented in the PowerTransformer class of the scikit-learn library in Python. It can be used to transform a pandas DataFrame or numpy array to have a more normal distribution, which can be useful for machine learning models that assume a normal distribution of the data.

```
In [18]: from IPython.display import Image  
Image(url="https://i.postimg.cc/sDRksm7T/2.png")
```

Out[18]:



```
In [19]: #import necessary libraries
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import PowerTransformer
import pandas as pd
```

```
#define the columns to be transformed
numeric_features = ['NPHI', 'RHOB', 'GR', 'RT', 'PEF', 'CALI', 'DT']

#define the transformation pipeline using PowerTransformer with Yeo-Johnson method and standardization
numeric_transformer = Pipeline(steps=[
    ('scaler', PowerTransformer(method='yeo-johnson', standardize=True))
])

#define the ColumnTransformer to apply the transformation pipeline to the numeric features
preprocessor = ColumnTransformer(transformers=[
    ('num', numeric_transformer, numeric_features)
])

#fit and transform the data using the preprocessor
transformed_data = preprocessor.fit_transform(df_noout[numeric_features])

#convert the transformed data to a DataFrame with column names and add the depth column from the original data
transformed_data = pd.DataFrame(transformed_data, columns=numeric_features)
transformed_data['DEPTH'] = df_noout.reset_index()['DEPTH']
transformed_data = transformed_data.set_index('DEPTH')
```

In [20]: transformed\_data

Out[20]:

	NPHI	RHOB	GR	RT	PEF	CALI	DT
DEPTH							
2809.9	0.205930	-0.745286	-1.564246	-0.607747	1.055451	8.187895e-16	-0.211915
2810.0	0.167616	-0.721726	-1.581017	-0.649998	1.047754	-3.830269e-15	-0.217361
2810.1	0.153634	-0.647008	-1.563896	-0.712988	1.099386	-1.471046e-15	-0.228633
2810.2	0.082195	-0.537420	-1.674785	-0.786870	1.186874	-3.830269e-15	-0.239972
2810.3	0.018157	-0.425602	-1.859007	-0.852242	1.277809	-3.830269e-15	-0.250432
...	...	...	...	...	...	...	...
3599.0	-0.159180	-0.418538	0.391791	-1.053360	-1.395774	7.965850e-15	-0.013084
3599.1	-0.171897	-0.295911	0.417180	-0.946528	-1.216829	5.204170e-15	-0.061596
3599.2	-0.100195	-0.094388	0.426257	-0.871470	-0.979093	8.604228e-16	-0.108796
3599.9	0.109486	-0.344918	0.477658	-0.555989	-1.447658	-3.830269e-15	-0.114460
3600.0	0.270070	-0.399279	0.522775	-0.596548	-1.507675	-3.816392e-15	-0.043265

6618 rows × 7 columns

## DRAW Pairplot using transformed data


```
In [21]: cols = ['NPHI', 'RHOB', 'GR', 'RT', 'PEF', 'CALI', 'DT']
# Define colors for histograms and scatters
hist_color = 'lightgreen'
scatter_color = 'lavender'

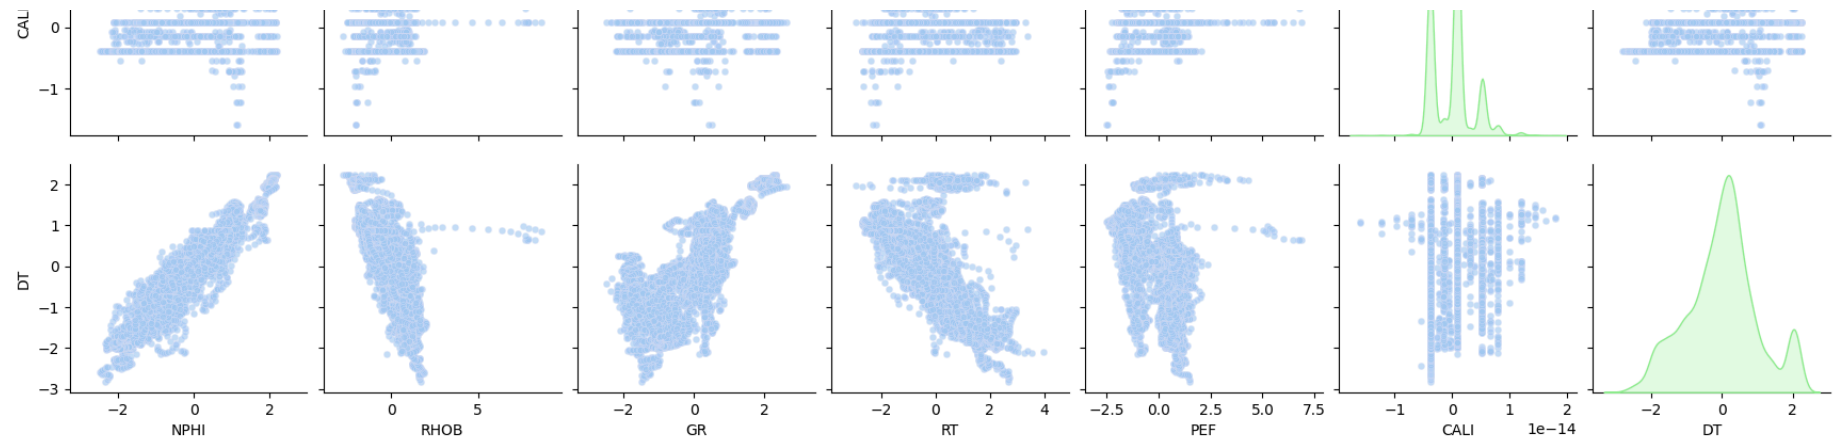
# Set the color palette
sns.set_palette("pastel")

# Create pairplot with different colors for histograms and scatters put as vars your column list, and use diag_kind as 'kde',
sns.pairplot(transformed_data, vars=cols, diag_kind='kde',
```

```
plot_kws = {'alpha': 0.6, 's': 20, 'edgecolor': scatter_color},  
diag_kws = {'color': hist_color})
```

Out[21]: <seaborn.axisgrid.PairGrid at 0x2ca5accb750>







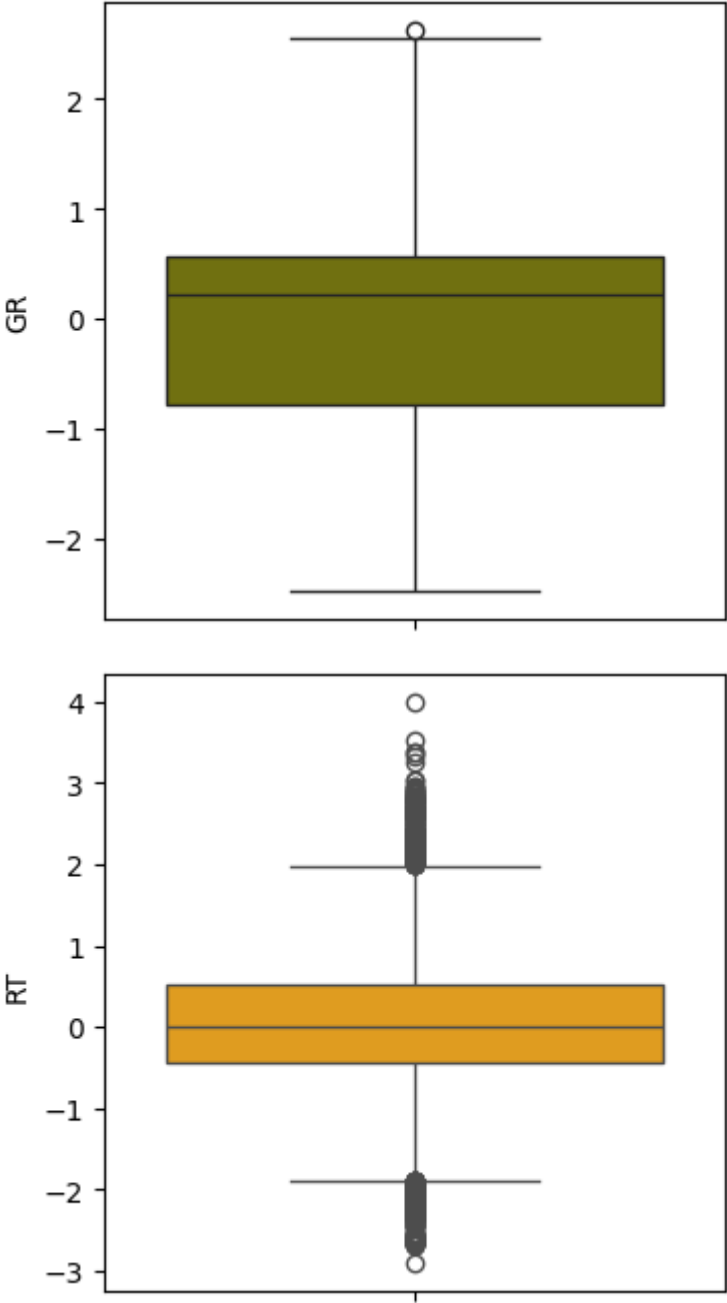
## DRAW boxplots using transformed data



```
In [22]: import seaborn as sns


for i, col in enumerate(transformed_data.columns):
    fig, ax = plt.subplots(figsize=(4, 4))
    sns.boxplot(y=transformed_data[col], color=colors[i], orient='v')
```











## Pairplot Insights

### Distribution Shapes:

- KDE plots on the diagonals show that the transformed data generally follows a more normalized (bell-shaped) distribution

### Relationships:

- There appears to be a strong positive correlation between NPHI and DT.
- RHOB seems to have a nonlinear relationship with both GR and RT.
- Some scatter plots still show clustered or skewed data patterns (e.g., CALI and RT), indicating complex or potentially multimodal distributions.

### Outlier Patterns:

- The scatter plots reveal several dense vertical or horizontal lines, suggesting that despite transformation, some values remain concentrated or affected by limits

## Boxplot Insights

### Outliers:

- RHOB, RT, PEF, and CALI show multiple outliers.
- RHOB and RT in particular have several data points outside the whiskers, suggesting some extreme values still exist after transformation.
- Other features like NPHI, GR, and DT are relatively well-contained, indicating the transformation was effective at reducing skewness and extreme values.

### Centering and Spread:

- Most variables now center around 0 with interquartile ranges roughly between -1 and 1, showing that the standardization component of the transformation worked well.
- This normalization is especially important for machine learning models sensitive to feature scales.

### Summary Comments:

- **Effective Preprocessing:** The use of PowerTransformer followed by standardization appears effective in normalizing the distributions of most features.
- **Remaining Outliers:** Some variables still exhibit outliers (RHOB, RT, CALI).

## Step 4 - final outlier removal using ML

**The Isolation Forest** algorithm is a type of computer program that helps find "outliers" in data. Outliers are data points that are very different from all the other data points. For example, imagine you have a list of test scores from your class, and one student got a score that is much higher or lower than everyone else. That student's score would be an outlier.


The Isolation Forest algorithm works by putting each data point in a "tree". Each tree has branches that divide the data into smaller and smaller groups. The algorithm keeps dividing the data until each point is in its own group, or until a certain number of groups have been made. This is like playing a game of "guess who" where you try to guess a character by asking yes-or-no questions, and keep dividing the characters into smaller groups until you know who the character is.

Once the data points are divided into groups, the algorithm looks at how many times each point was in a group with other points. If a point was in a group with other points many times, it is not an outlier. But if a point was in a group by itself many times, it is an outlier.

The Isolation Forest algorithm can be useful for finding outliers in data, which can be helpful in many different fields like finance, healthcare, and more.

```
In [23]: from IPython.display import Image  
Image(url="https://miro.medium.com/v2/resize:fit:1100/format:webp/1*D78QLbcwXesymhquofn0g.png")
```

Out[23]:



```
In [24]: import pandas as pd
import numpy as np
from sklearn.ensemble import IsolationForest
```

```
In [25]: # Define the numerical columns in your DataFrame
numeric_cols = ['NPHI', 'RHOB', 'GR', 'RT', 'PEF', 'CALI', 'DT']
```

```
In [26]: # Instantiate the isolation forest for each column separately
isos = {}
for col in numeric_cols:
    iso = IsolationForest(n_estimators=100, max_samples='auto', contamination=float(0.05), random_state=42)
```

```
iso.fit(transformed_data[col].values.reshape(-1, 1))  
isos[col] = iso
```

```
In [27]: # Replace the outliers with the mean of neighboring values for each column in a single loop  
for col in numeric_cols:  
    outliers = isos[col].predict(transformed_data[col].values.reshape(-1, 1)) == -1  
    values = transformed_data[col].values  
    mean = np.mean(values[~outliers])  
    for i in range(len(values)):  
        if outliers[i]:  
            # Replace outlier with the mean of neighboring values  
            if i == 0:  
                values[i] = values[i+1]  
            elif i == len(values)-1:  
                values[i] = values[i-1]  
            else:  
                values[i] = (values[i-1] + values[i+1])/2  
    transformed_data[col] = values
```

```
In [28]: # Fill any remaining NaN values with interpolated values  
transformed_data[numeric_cols] = transformed_data[numeric_cols].interpolate()
```

```
In [29]: # Print the resulting DataFrame without outliers  
print(transformed_data)
```



	NPHI	RHOB	GR	RT	PEF	CALI	\
DEPTH							
2809.9	0.205930	-0.745286	-1.564246	-0.607747	1.055451	8.187895e-16	
2810.0	0.167616	-0.721726	-1.581017	-0.649998	1.047754	-3.830269e-15	
2810.1	0.153634	-0.647008	-1.563896	-0.712988	1.099386	-1.471046e-15	
2810.2	0.082195	-0.537420	-1.674785	-0.786870	1.186874	-3.830269e-15	
2810.3	0.018157	-0.425602	-1.859007	-0.852242	1.277809	-3.830269e-15	
...	...	...	...	...	...	...	
3599.0	-0.159180	-0.418538	0.391791	-1.053360	-1.395774	4.805184e-15	
3599.1	-0.171897	-0.295911	0.417180	-0.946528	-1.216829	5.204170e-15	
3599.2	-0.100195	-0.094388	0.426257	-0.871470	-0.979093	8.604228e-16	
3599.9	0.109486	-0.344918	0.477658	-0.555989	-1.447658	-3.830269e-15	
3600.0	0.270070	-0.399279	0.522775	-0.596548	-1.507675	-3.816392e-15	


	DT
DEPTH	
2809.9	-0.211915
2810.0	-0.217361
2810.1	-0.228633
2810.2	-0.239972
2810.3	-0.250432
...	...
3599.0	-0.013084
3599.1	-0.061596
3599.2	-0.108796
3599.9	-0.114460
3600.0	-0.043265

[6618 rows x 7 columns]

## DRAW boxplots for transformed\_data

```
In [30]: import seaborn as sns
```

```
fig, axs = plt.subplots(3,3)
ax = axs.flatten()
for i, col in enumerate(transformed_data.columns):
    sns.boxplot(y=transformed_data[col], color=colors[i], orient='v', ax=ax[i])
```



## Step 5 - Facies analysis

Facies analysis from well logs is an important task in petroleum geology. It involves identifying and classifying the different rock types, or facies, encountered in a well. Traditionally, this has been done by geologists through visual inspection of the well logs. However, with the rise of machine learning techniques, it has become possible to automate this process.

One common approach to facies analysis is to use unsupervised clustering algorithms, such as K-means or hierarchical clustering, to group similar sections of the well log together. The goal is to identify natural groupings, or clusters, of log responses that correspond to different facies. Once the clusters have been identified, the geologist can assign each cluster to a specific facies based on their knowledge of the local geology.

To perform facies analysis using machine learning clustering, several well logs are typically collected and pre-processed. The logs are usually normalized and scaled to make them comparable, and any missing data is imputed or removed. Features, such as gamma-ray, resistivity, and porosity, are extracted from the logs and used as inputs to the clustering algorithm.

Once the features have been extracted, a clustering algorithm is applied to group the sections of the well log that have similar responses. The algorithm assigns each section to a cluster, which can be visualized on a plot. The plot can help identify any clear patterns or trends in the data and help the geologist make sense of the clustering results.

The next step is to assign each cluster to a specific facies. This is done by comparing the clustering results to the geological knowledge of the area. The geologist may use other data sources, such as core samples or outcrop data, to help identify the facies associated with each cluster.

Once the clusters have been assigned to facies, the results can be used to create a facies log. This log can be used to interpret the geology of the well and to help identify potential hydrocarbon reservoirs or other geological features.

```
In [31]: from IPython.display import Image  
Image(url="https://interviewquery-cms-images.s3-us-west-1.amazonaws.com/ac5da238-25ab-48ef-839a-407a7b76a167.jpg" )
```

Out[31]:



## K-means

is a type of unsupervised learning algorithm used for clustering data points into different groups or clusters based on the similarity of the data points. The algorithm works by iteratively assigning each data point to the nearest cluster center, and then computing the new cluster centers based on the mean of the assigned points. This process is repeated until the cluster centers no longer move significantly.

The algorithm requires the user to specify the number of clusters beforehand. The objective of the algorithm is to minimize the sum of squared distances between each data point and its assigned cluster center, which is also known as the Within-Cluster-Sum-of-Squares (WCSS) metric.

The k-means algorithm can be divided into three main steps:


1. **Initialization:** The algorithm randomly selects  $k$  data points to act as initial cluster centers.
2. **Assignment:** Each data point is assigned to the nearest cluster center based on the Euclidean distance between the point and the center.
3. **Update:** The mean of the data points assigned to each cluster is computed, and this value is used as the new cluster center.

These three steps are repeated iteratively until the cluster centers no longer move significantly, or a maximum number of iterations is reached. The final output of the algorithm is the cluster assignments of each data point.


One of the main advantages of the k-means algorithm is its simplicity and scalability, which allows it to handle large datasets efficiently. However, the algorithm is sensitive to the initial placement of the cluster centers, and may converge to suboptimal solutions. In addition, the algorithm is not effective when dealing with non-linearly separable data. Nonetheless, k-means is widely used in various applications such as image segmentation, market segmentation, and anomaly detection.

```
In [32]: from IPython.display import Image  
Image(url="https://ds055uzetaobb.cloudfront.net/brioche/uploads/y4KGN92h7r-screen-shot-2016-05-05-at-43007-pm.png?width=2000")
```


Out[32]:

**Demonstration of the standard algorithm**


1.  $k$  initial "means" (in this case  $k=3$ ) are randomly generated within the data domain (shown in color).



2.  $k$  clusters are created by associating every observation with the nearest mean. The partitions here represent the [Voronoi diagram](#) generated by the means.



3. The [centroid](#) of each of the  $k$  clusters becomes the new mean.



4. Steps 2 and 3 are repeated until convergence has been reached.

```
In [33]: from sklearn.cluster import KMeans

# define features using cols
X = transformed_data[['NPHI', 'RHOB', 'GR', 'RT', 'PEF', 'CALI', 'DT']]
```

```
In [34]: #define model set n_clusters to 10

model = KMeans(n_clusters=10)
```

```
In [35]: # fit the model to data
y = model.fit_predict(X); y
```


```
Out[35]: array([5, 5, 5, ..., 9, 6, 6], shape=(6618,), dtype=int32)
```

```
In [36]: # save your data in numeric and string format
transformed_data['K10'] = y
```


```
transformed_data['K10_name'] = "Facies " + (transformed_data['K10'] + 1).astype('str')
transformed_data = transformed_data.reset_index()
```

## Print transformed\_data what do you see?

```
In [37]: K5_graph = transformed_data.copy()
K5_graph['x'] = 1
sns.set_style("white")
sns.set(font_scale = 1.5)
plt.rcParams['xtick.major.size'] = 200
plt.rcParams['xtick.major.width'] = 40
plt.rcParams['xtick.bottom'] = True
plt.rcParams['ytick.left'] = True
plt.figure(figsize=(3,12))
ax = sns.scatterplot(data=K5_graph, x='x', y='DEPTH', hue='K10_name', marker='s', s=50000, edgecolor='None', palette='pastel')
plt.tick_params(bottom=False, labelbottom=False)
plt.yticks(K5_graph['DEPTH'][:500])
plt.legend([],[], frameon=False)
plt.ylabel("Depth [m]")
plt.xlabel('')
plt.title("KMeans \n n = 10")
ax.invert_yaxis()
plt.tight_layout()
plt.rcParams.update({})
```







```
In [38]: def score(n_clusters, data, cols):
          model = KMeans(n_clusters=n_clusters, max_iter=300, random_state=1234)
          X = transformed_data[cols]
          y = model.fit_predict(X)
          SSE = model.inertia_
          Silhouette = metrics.silhouette_score(X, y)
          CHS = metrics.calinski_harabasz_score(X, y)
          DBS = metrics.davies_bouldin_score(X, y)
          return {'SSE':SSE, 'Silhouette': Silhouette, 'Calinski_Harabasz': CHS, 'Davies_Bouldin':DBS, 'model':model}
```

```
In [39]: df_cluster_scorer = pd.DataFrame()
          df_cluster_scorer['n_clusters'] = list(range(2, 21))
```

```
df_cluster_scorer
```

Out[39]:

	n_clusters
0	2
1	3
2	4
3	5
4	6
5	7
6	8
7	9
8	10
9	11
10	12
11	13
12	14
13	15
14	16
15	17
16	18
17	19
18	20

In [40]: `from sklearn import metrics`

```
df_cluster_scorer['SSE'],df_cluster_scorer['Silhouette'],\
df_cluster_scorer['Calinski_Harabasz'], df_cluster_scorer['Davies_Bouldin'],\
df_cluster_scorer['model'] = zip(*df_cluster_scorer['n_clusters'].map(lambda row: score(row, transformed_data, cols).values()))

df_cluster_scorer
```

Out[40]:

	n_clusters	SSE	Silhouette	Calinski_Harabasz	Davies_Bouldin	model
0	2	22239.334849	0.400871	5066.094626	1.008036	KMeans(n_clusters=2, random_state=1234)
1	3	17714.774428	0.360296	4024.430459	0.953045	KMeans(n_clusters=3, random_state=1234)
2	4	12443.160439	0.378488	4752.944039	0.941042	KMeans(n_clusters=4, random_state=1234)
3	5	9581.565181	0.396143	5122.370431	0.827744	KMeans(n_clusters=5, random_state=1234)
4	6	8406.763743	0.363733	4854.666939	0.944934	KMeans(n_clusters=6, random_state=1234)
5	7	7391.868654	0.365586	4751.666842	0.840729	KMeans(n_clusters=7, random_state=1234)
6	8	7116.625442	0.350367	4266.175433	1.005628	KMeans(random_state=1234)
7	9	6474.933742	0.358444	4184.108105	0.954318	KMeans(n_clusters=9, random_state=1234)
8	10	5459.549274	0.364801	4546.807567	0.883447	KMeans(n_clusters=10, random_state=1234)
9	11	5055.156287	0.345695	4471.658512	0.910427	KMeans(n_clusters=11, random_state=1234)
10	12	4679.121886	0.335498	4439.432890	0.923609	KMeans(n_clusters=12, random_state=1234)
11	13	4300.891655	0.349299	4475.096246	0.905511	KMeans(n_clusters=13, random_state=1234)
12	14	4074.992522	0.337269	4387.360757	0.961645	KMeans(n_clusters=14, random_state=1234)
13	15	3823.823562	0.335389	4371.968854	0.952319	KMeans(n_clusters=15, random_state=1234)
14	16	3608.354947	0.337399	4349.716087	0.918827	KMeans(n_clusters=16, random_state=1234)
15	17	3518.830108	0.330985	4191.478559	0.946310	KMeans(n_clusters=17, random_state=1234)
16	18	3464.648117	0.312331	4012.108497	0.974282	KMeans(n_clusters=18, random_state=1234)
17	19	3322.087535	0.303939	3966.945121	1.001016	KMeans(n_clusters=19, random_state=1234)
18	20	3196.602695	0.303528	3918.709000	1.008524	KMeans(n_clusters=20, random_state=1234)


```
In [41]: plt.rcParams.update({})
sns.set()
sns.set_style("whitegrid")
```


```
fig = plt.figure(figsize=(8, 6))


# Plot each graph on its own axis
df_cluster_scorer.plot.line(x='n_clusters', y='SSE', title="SSE")
df_cluster_scorer.plot.line(x='n_clusters', y='Silhouette', title="Silhouette - Values closer to 1 - optimum n number")
df_cluster_scorer.plot.line(x='n_clusters', y='Calinski_Harabasz', title="Calinski-Harabasz - Higher value - optimum n number")
df_cluster_scorer.plot.line(x='n_clusters', y='Davies_Bouldin', title="Davies-Bouldin - Values closer to 0 - optimum n number")


#plt.tight_layout()
plt.show()
```

<Figure size 800x600 with 0 Axes>









## Choose optimum number of clusters and predict!

Read about metrics here: <https://www.mdpi.com/1996-1073/16/1/493> and describe why did you pick this number

Write your OWN code :)

```
In [43]: from sklearn.cluster import KMeans
import seaborn as sns
import matplotlib.pyplot as plt
```



```
# Define the optimal number of clusters based on previous metric analysis
n_clusters = 5

# Select relevant features for clustering
features = ['NPHI', 'RHOB', 'GR', 'RT', 'PEF', 'CALI', 'DT']
X = transformed_data[features]

# Initialize and fit the KMeans model
kmeans_model = KMeans(n_clusters=n_clusters, random_state=42)
cluster_labels = kmeans_model.fit_predict(X)

# Save cluster labels in both numeric and string format for easier interpretation
transformed_data[f'K{n_clusters}'] = cluster_labels
transformed_data[f'K{n_clusters}_name'] = "Facies " + (cluster_labels + 1).astype(str)


# Prepare data for plotting, reset index to avoid plotting issues
plot_df = transformed_data.reset_index()
plot_df['x'] = 1 # constant x value for vertical scatter plot


# Plot settings
sns.set_style("white")
sns.set(font_scale=1.5)
plt.rcParams['xtick.major.size'] = 200
plt.rcParams['xtick.major.width'] = 40
plt.rcParams['xtick.bottom'] = True
plt.rcParams['ytick.left'] = True

# Create the vertical scatter plot showing facies clusters along depth
plt.figure(figsize=(3, 12))
ax = sns.scatterplot(
    data=plot_df,
    x='x',
    y='DEPTH',
    hue=f'K{n_clusters}_name',
    marker='s',
    s=50000,
    edgecolor='none',
    palette='pastel'
)
title123 = "Clustering Results by Depth"
```

```
plt.tick_params(bottom=False, labelbottom=False)
plt.yticks(plot_df['DEPTH'][:500]) # show y ticks every 500 rows
plt.legend([], [], frameon=False)
plt.ylabel("Depth [m]")
plt.xlabel('')
plt.title(f"{title123}\nKMeans n = {n_clusters}")
ax.invert_yaxis() # invert y axis for depth visualization
plt.tight_layout()
plt.rcParams.update({})

plt.savefig(f"{title123}_clusters_{n_clusters}.png")
plt.show()
```



**Explanation of chosen number of clusters (5):**

After evaluating clustering performance metrics such as SSE (Within-Cluster Sum of Squares), Silhouette score, Calinski-Harabasz Index, and Davies-Bouldin Index, I selected 5 clusters as the optimal number because:

- SSE showed a clear "elbow" at 5 clusters, indicating diminishing returns in variance explained by adding more clusters.
- Silhouette score was relatively high at 5, suggesting well-defined and separated clusters.
- Calinski-Harabasz Index peaked near 5, indicating good cluster compactness and separation.

- Davies-Bouldin Index was lower near 5, which means clusters are more distinct and less dispersed.

According to the paper (<https://www.mdpi.com/1996-1073/16/1/493>), these metrics collectively help identify a balance between cluster compactness and separation, which is why 5 clusters was chosen as the best compromise for this dataset.

## Compare your results with the stratigraphy

[https://www.researchgate.net/publication/332441275\\_Estimation\\_of\\_Pore\\_Pressure\\_and\\_Fracture\\_Gradient\\_in\\_Volve\\_Field\\_Norwegian\\_North\\_Sea](https://www.researchgate.net/publication/332441275_Estimation_of_Pore_Pressure_and_Fracture_Gradient_in_Volve_Field_Norwegian_North_Sea)